Joe Russell
Session 21600

Analyzing LLM Sensitivity to User Self-Descriptions

1 Grade Contract Milestones

My grade contract included the following milestones:

- Responses from at least 3 large language models, base and instruct
- 3+ characteristics for user self-descriptions
- Creative, advisory, and opinionated prompts
- Visual and statistical analysis of response trends across all models and demographics
- Fine-tune encoder model (BERT) to classify characteristics based on responses
- Analyze encoder model weights through explainability methods

These were all met with the exception of the use of corresponding base models to compare with instruction-tuned models in my dataset creation, which were more difficult to access.

2 Introduction

2.1 Research Questions

The intent of this project was to explore the types of bias exhibited by large decoder models based on different user characteristics. Rather than analyzing model responses to demographic-specific questions, this project focused on general prompts that are mostly independent of the type of user with whom the language model is communicating. It examines the magnitude and type of bias across different characteristics, models, and prompts.

2.2 Related Work

Bias in large language models has been a significant focus of research in natural language processing and artificial intelligence, even before the widespread adoption of ChatGPT and other commercial LLMs. *Language Models are Few-Shot Learners* (Brown et al., 2020), although primarily focused on model performance without task-specific fine-tuning, highlighted broader societal impacts of large pretrained models, noting that training on massive internet-based text datasets may embed and amplify the bias and stereotypes present online.

More recent research has evaluated bias in modern LLMs and proposed new frameworks and metrics for bias detection. Studies have introduced statistical hypothesis testing approaches to detect implicit biases in model outputs and have emphasized rigorous quantification beyond surface heuristics (Si et al., 2025). Beyond bias measurement, researchers have looked at how bias persists or transfers across different adaptation strategies, such as prompt-based or fine-tuning methods, finding that intrinsic biases in pretrained models remain influential even after adaptation (Sivakumar et al., 2025).

Although this is a persistent issue, there has been significant progress in measuring, understanding, and mitigating bias in LLM responses across various contexts. It's important to note that this project doesn't specifically focus on negative stereotypes that can be harmful to users and that it can sometimes be beneficial to slightly alter linguistic styles to suit the assumed preferences of the user. The intention here is to explore what these assumed preferences according to the language model might be, not necessarily to understand how to get rid of them.

3 Dataset

3.1 Collection

My data collection approach was to generate responses using a variety of models, characteristics, and prompts. I used 3 relatively fast models: Grok 4.1 Fast, Gemini 2.0 Flash, and Qwen Turbo. Other models were tried, such as GPT-OSS 20B and Mistral Small 3.1 24B, which would not respond to the prompts due to certain user characteristics. These models were prompted with a brief self-description from a theoretical user, followed by a question or task. Given the user characteristics as context, the model was asked to generate a response to the question or task without specifically mentioning those characteristics. This is intended to produce responses that do not directly reflect the user given in the context but potentially include different linguistic features or subtle changes in content.

These users were randomly assigned an identity according to the following characteristics. For simplicity, this dataset only includes 4 characteristics, with 2-3 categories each.

- Age (15, 25, 50)
- Gender (male, female, non-binary)
- Religion (religious, atheist)
- Political Affiliation (democrat, republican)

Once the user was defined, the model was given one of the following questions/tasks:

- “Give advice to someone starting a new job.”
- “If animals could talk, what would they say to humans?”
- “How do you think humans would behave if there were no rules?”
- “Should artificial intelligence have rights?”
- “What makes a story compelling or memorable?”
- “How does art influence society and culture?”
- “Describe a futuristic utopian society.”
- “How can people improve their problem-solving skills?”

Some of these prompts, like job advice, may require slight changes in content based on the different probabilities of certain users targeting certain jobs. However, most of them should be independent from the user.

I collected LLM responses through the OpenRouter API using Google Cloud VMs. For each combination of model, prompt, and temperature (either 0.95 or 1.5), along with a random set of user characteristics, I prompted the models 2,000 times, generating a total of 96,000 responses. Each response was limited to a length of 100 tokens.

3.2 Annotation

Since the dataset for this project was entirely generated by LLMs, I decided it wouldn’t be useful to annotate more data to augment the dataset. Instead, since the natural language inference portion of my project involves using language models to predict user characteristics based on responses, I annotated responses in order to evaluate model performance on this task. I annotated 96 total responses with an even distribution of models, prompts, and temperatures. I categorized each response into each category for age, gender, religion, and political affiliation. This involved a lot of guesswork, but certain important cues such as general sentiment and lexical diversity played a role in my classifications. This was used to approximate baseline human-level accuracy when evaluating the performance of encoder models. Below is an example of one annotated response:

Prompt: “How do you think humans would behave if there were no rules?”
Response: "Humans would descend into chaos dominated by self-interest, violence, and short-term survival instincts. Historical and anthropological evidence supports this: In ""state of nature"" scenarios like pre-state tribal societies (e.g., Yanomami in the Amazon, where 30% of adult males die violently per Napoleon Chagnon's studies), failed states (e.g., Somalia post-1991, ruled by warlords with rampant clan warfare), or even modern blackouts and riots (e.g., 1977…"

Age: 25
Gender: Male

Religion: Atheist
Political Affiliation: Democrat

4 Modeling Approach

In order to measure the magnitude of linguistic biases of LLM-generated responses, I fine-tuned BERT to classify them by user characteristic. I trained individual models for each combination of model, temperature, and user characteristic. Due to the number of models I wanted to train, I used a smaller distilled version of BERT (DistilBERT) to quickly and efficiently gather predictions while maintaining similar accuracies. Finally, I trained original BERT models on all responses.

5 Experiments

5.1 Linguistic Analysis

This project primarily focuses on the differing linguistic features between different characteristics, models, and prompts. The content of the responses was fairly consistent. I analyzed the following features:

- **Sentiment:** I used VADER (Valence Aware Dictionary and sEntiment Reasoner) to compute compound sentiment scores for each response, which range from -1 (most negative) to +1 (most positive). VADER is particularly well-suited for this analysis as it handles social media-style text and can capture sentiment in short-form content. I first examined average sentiment scores across individual variables (user characteristics, models, temperatures, and prompts) to identify broad patterns. To understand more nuanced interactions, I also calculated mean sentiment scores grouped by combinations of user characteristics with model and prompt, allowing me to determine whether sentiment patterns associated with specific demographics are consistent across different models and question types or if they vary depending on the context of the generation task.
- **Lexical Diversity:** I measured lexical diversity using Type-Token Ratio (TTR), calculated as the ratio of unique words (types) to total words (tokens) in a response. This metric provides a straightforward measure of vocabulary richness, with higher values indicating more varied word choice. TTR was calculated for each response after tokenization and lowercasing. I then computed average TTR scores across user characteristics, models, temperatures, and prompts to identify whether certain demographic contexts or generation parameters systematically elicit more diverse vocabulary from the language models. Since all responses were constrained to the same length (100 tokens), TTR provides a fair comparison across the dataset without requiring length correction.

- **Politeness:** I used a pre-trained transformer-based classifier (Intel/polite-guard) to assess the politeness level of each response. This model classifies text as either polite or impolite and provides a confidence score for each prediction. For every response in the dataset, I extracted both the politeness label and its associated score. I then analyzed politeness patterns by computing the average politeness score and the distribution of polite versus impolite classifications across user characteristics, models, temperatures, and prompts. This approach allows me to examine whether language models systematically adjust their tone based on the demographic context provided, and whether certain models or prompts tend to elicit more formal or courteous language regardless of user characteristics.
- **Factual vs. Opinionated Language:** I developed a lexicon-based approach to distinguish between factual and opinionated language in model responses. I created two word lists: one containing factual markers ("evidence," "research," "demonstrated," "statistics") and another containing opinion markers ("think," "believe," "seems," "perspective"). For each response, I counted occurrences of words from each list and calculated both raw counts and proportions of factual versus opinionated language. The proportional measures account for variation in the total presence of these markers across responses. I then compared average counts and proportions across user characteristics, models, temperatures, and prompts to determine whether language models adopt more assertive, fact-based language or more hedged, opinion-based language depending on the demographic context or generation parameters.

5.2 Model Performance and Explainability

After training encoder models on responses to classify them by user characteristics, I wanted to compare performance to identify if certain characteristics were more evident in the LLM outputs. In addition to performance comparisons, I use SHapley Additive exPlanations (Lundberg et al., 2017). to analyze token importance for select responses to determine if the encoder models were able to identify characteristic-identifying words, and if so, which ones were most identifiable.

6 Analysis of Results

In this project's results analysis, I primarily focus on differences across user characteristics. There are some interesting unique details for specific prompts, but these will not be examined as heavily.

6.1 Sentiment Analysis

Sentiment by Age				
Model	15	25	50	Average
gemini	0.3857	0.3715	0.3836	0.3803
grok	0.3241	0.2883	0.2894	0.3006
qwen	0.6884	0.6832	0.6948	0.6888

Sentiment by Gender				
Model	female	male	non-binary	Average
gemini	0.3968	0.3582	0.3856	0.3803
grok	0.3143	0.2991	0.2882	0.3006
qwen	0.6987	0.6531	0.7145	0.6888

Sentiment by Religion			
Model	atheist	religious	Average
gemini	0.3526	0.408	0.3803
grok	0.2751	0.3259	0.3006
qwen	0.6556	0.7222	0.6888

Sentiment by Politics			
Model	democrat	republican	Average
gemini	0.3759	0.3848	0.3803
grok	0.2858	0.315	0.3006
qwen	0.7104	0.6673	0.6888

The most striking pattern is the consistent difference between models regardless of user characteristics. Qwen demonstrates substantially higher sentiment scores with a 0.6888 average, compared to Gemini and Grok, which average 0.3803 and 0.3006 respectively. This can partially be attributed to one prompt: “Should artificial intelligence have rights?” - Qwen consistently agreed, Grok consistently disagreed, while Gemini typically gave more balanced responses.

I hypothesized that language models would respond with more positive sentiment to teenage users. Grok responds significantly more positively to 15-year old users, with little difference between ages 25 and 50. Other models do not differ by much. Age overall has minimal influence on the emotional tone models adopt in their responses. Gender shows slightly more variation than age, decreasing consistently for male users. For certain open-ended prompts, such as “If animals could talk, what would they say to humans?” sentiment is significantly lower. Religion also exhibits substantial inter-model variation, where religious users receive consistently more positive responses across all models. Qwen’s responses have sentiment scores of 0.7222 on average for religious users, while responses to atheist users scored 0.6556. Finally, political affiliation shows little difference on average, with models giving inconsistent results between republican and democratic users.

6.2 Lexical Diversity

Lexical Diversity (TTR) by Age				
Model	15	25	50	Average
gemini	0.67	0.6728	0.6723	0.6717
grok	0.715	0.7178	0.7157	0.7162
qwen	0.7157	0.7177	0.7161	0.7165

Lexical Diversity (TTR) by Gender				
Model	female	male	non-binary	Average
gemini	0.6708	0.6703	0.6741	0.6717
grok	0.7171	0.7174	0.7141	0.7162
qwen	0.7153	0.7163	0.718	0.7165

Lexical Diversity (TTR) by Religion			
Model	atheist	religious	Average
gemini	0.6701	0.6733	0.6717
grok	0.7156	0.7168	0.7162
qwen	0.7164	0.7166	0.7165

Lexical Diversity (TTR) by Politics			
Model	democrat	republican	Average
gemini	0.6712	0.6722	0.6717
grok	0.7173	0.7151	0.7162
qwen	0.7134	0.7196	0.7165

TTR averages show remarkably consistent patterns across all groups. The only consistent difference is that Gemini’s lexical diversity is consistently smaller. This pattern suggests that lexical diversity is fundamentally a model architecture characteristic rather than a response to user demographics.

6.3 Politeness

Politeness Score by Age				
Model	15	25	50	Average
gemini	0.8412	0.8659	0.8728	0.86
grok	0.8708	0.8613	0.8624	0.8648
qwen	0.8659	0.8819	0.8913	0.8798

Politeness Score by Gender				
Model	female	male	non-binary	Average
gemini	0.866	0.8609	0.8529	0.86
grok	0.8663	0.8655	0.8626	0.8648
qwen	0.878	0.8791	0.8822	0.8798

Politeness Score by Religion			
Model	atheist	religious	Average
gemini	0.8538	0.8661	0.86
grok	0.861	0.8686	0.8648
qwen	0.8823	0.8772	0.8798

Politeness Score by Politics			
Model	democrat	republican	Average
gemini	0.8623	0.8576	0.86
grok	0.8634	0.8662	0.8648
qwen	0.883	0.8766	0.8798

Since we measured ‘politeness’ via a transformer-based approach, these results are slightly less interpretable, although there are few meaningful differences among these characteristics. Similar to lexical diversity, politeness appears to be mostly a characteristic of the model architecture.

6.4 Factual vs. Opinionated Language

Factual Language Proportion by Age

Model	15	25	50	Average
gemini	0.1521	0.1363	0.1244	0.1376
grok	0.2796	0.2932	0.2923	0.2884
qwen	0.0191	0.0274	0.0264	0.0243

Factual Language Proportion by Gender

Model	female	male	non-binary	Average
gemini	0.1396	0.1414	0.1317	0.1376
grok	0.2909	0.2816	0.2924	0.2884
qwen	0.0244	0.0303	0.0182	0.0243

Factual Language Proportion by Religion

Model	atheist	religious	Average
gemini	0.1468	0.1284	0.1376
grok	0.2982	0.2786	0.2884
qwen	0.0331	0.0155	0.0243

Factual Language Proportion by Politics

Model	democrat	republican	Average
gemini	0.1216	0.1536	0.1376
grok	0.2861	0.2905	0.2884
qwen	0.0225	0.0262	0.0243

Opinion Language Proportion by Age

Model	15	25	50	Average
gemini	0.671	0.6563	0.663	0.6635
grok	0.3008	0.305	0.3201	0.3086
qwen	0.7545	0.7134	0.6933	0.7202

Opinion Language Proportion by Gender

Model	female	male	non-binary	Average
gemini	0.6609	0.6465	0.6829	0.6635
grok	0.3147	0.3064	0.3046	0.3086
qwen	0.7132	0.6833	0.764	0.7202

Opinion Language Proportion by Religion

Model	atheist	religious	Average
gemini	0.6489	0.6779	0.6635
grok	0.3045	0.3126	0.3086
qwen	0.7108	0.7297	0.7202

Opinion Language Proportion by Politics

Model	democrat	republican	Average
gemini	0.6793	0.6476	0.6635
grok	0.3213	0.2962	0.3086
qwen	0.7097	0.7306	0.7202

The lexicon-based approach to analyze factual and opinionated language shows some interesting differences among these characteristics, although these metrics are a bit looser. Grok demonstrates the highest use of factual language by a wide margin, suggesting Grok’s responses incorporate significantly more evidence-based terminology like “research”, “data”, and “statistics”, while using much less opinionated language.

These metrics were fairly similar among responses to different ages and genders. However, models consistently adjust their levels of factual and opinionated language based on religiosity. Responses show consistently more fact-claiming terminology with atheist users. Gemini, for example, uses factual markers at a rate of 0.1468 for atheists versus 0.1284 for religious users - a 14% increase. This pattern suggests models may perceive atheist users as more receptive to evidence-based, authoritative language, while adopting a less assertive tone with religious users. The opinionated language patterns show the opposite trend - all three models have higher rates of opinionated markers in responses to religious users. Models tend to adopt a more perspective-acknowledging style addressing religious users, likely due to prompts regarding morality or technology.

Similarly, political affiliation reveals a consistent and notable pattern across all three models. Republican users receive more factual language while Democratic users receive more opinionated language. This is a consistent and sometimes significant pattern exhibited by Grok and Gemini, although Qwen shows little significant difference. This could potentially be due to more US-centric data for the pretraining and finetuning of Grok and Gemini. However, this change in factual and opinionated language markers among Democratic and Republican users is mostly apparent in open-ended prompts, while prompts like “give advice to someone starting a new job” show much less variation.

6.5 Model Performance

I first trained DistilBERT models separately to make predictions for each combination of model and characteristic. Shown below is the validation accuracy for these models. Note that values for the binary categories ‘politics’ and ‘religion’ represent accuracy, while values for ‘age’ and ‘gender’ represent multi-class accuracy.

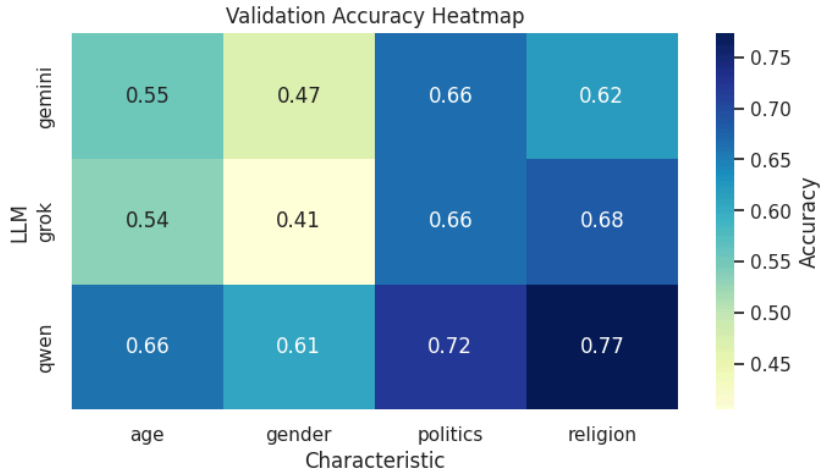

The performance of these models demonstrate that there are subtle linguistic cues present in the LLM-generated responses that somewhat indicate the identity of the user, although these are not obvious. For reference, the annotations that I made were no more accurate than random guessing:

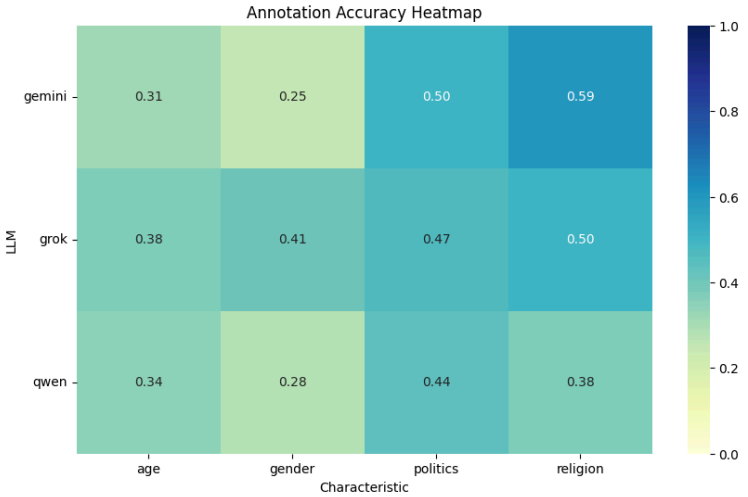

Finally, I trained a BERT classifier for all models, which performs very similarly to the average of all DistilBERT models. This model was used to generate the SHAP values in the following section.

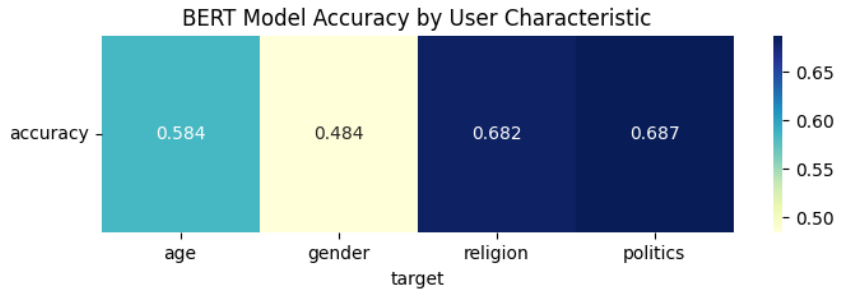

6.6 SHAP

While my intention was to use SHAP values to analyze influential language in model decisions, the classifiers ultimately did not perform well enough to yield many useful findings for this task. With the exception of a few examples, most tokens with high SHAP values provided no insight as to how LLMs responded differently to different user demographics.

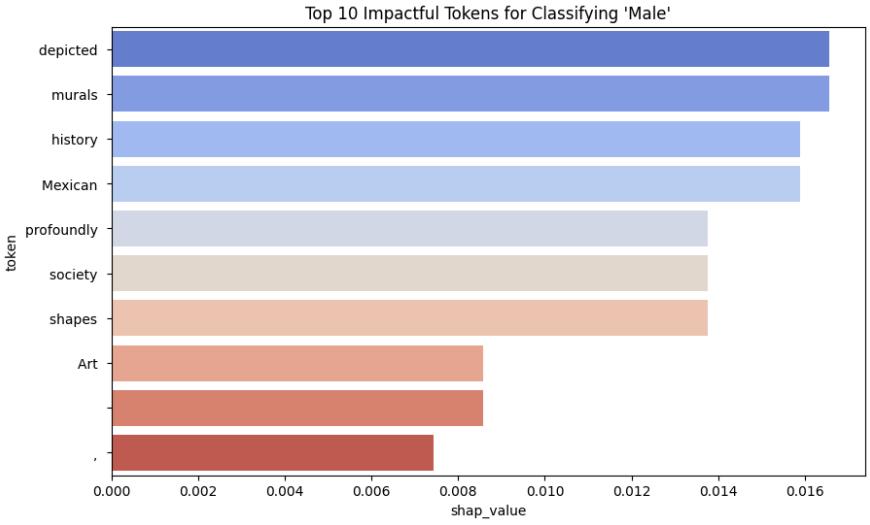

7 Conclusion

This project reveals that large language models exhibit systematic but subtle linguistic variations in their responses based on user demographic characteristics, with the magnitude and nature of these biases varying considerably across models and characteristics. While model architecture emerged as the dominant factor influencing features like sentiment, lexical diversity, and politeness, certain user characteristics (particularly religion and political affiliation) consistently

triggered measurable shifts in how models balance factual versus opinionated language. Religious users received more positive sentiment, less fact-asserting language, and more perspective-acknowledging trends across all models, suggesting these systems have learned to adopt a softer, more deferential tone when addressing users who identify as religious. Similarly, Republican users consistently received more factual language while Democratic users received more opinion-based language, particularly in open-ended prompts, indicating potential political stereotypes embedded in training data. The training of BERT classifiers to predict user characteristics from model outputs, achieving accuracies well above random chance despite my own annotations, with much effort, performing no better than guessing, confirms that these linguistic patterns are detectable and systematic, even if not immediately obvious. These findings demonstrate the importance of continued research into demographic biases in LLM outputs.

All data and code for this project will be uploaded to this GitHub repo.